

ASSIGNMENT 6:

Implement K-Means clustering/ hierarchical clustering on sales_data_sample.csv dataset. Determine the number of clusters using the elbow method. Dataset link : <https://www.kaggle.com/datasets/kyanyoga/sample-sales-data>

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
#Importing the required libraries.
```

```
In [2]: from sklearn.cluster import KMeans, k_means #For clustering
from sklearn.decomposition import PCA #Linear Dimensionality reduction.
```

```
In [3]: df = pd.read_csv("sales_data_sample.csv") #Loading the dataset.
```

Preprocessing

```
In [4]: df.head()
```

```
Out[4]:
```

	ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	ORDERDATE	STATUS	QTR_ID
0	10107	30	95.70	2	2871.00	2/24/2003 0:00	Shipped	1
1	10121	34	81.35	5	2765.90	5/7/2003 0:00	Shipped	2
2	10134	41	94.74	2	3884.34	7/1/2003 0:00	Shipped	3
3	10145	45	83.26	6	3746.70	8/25/2003 0:00	Shipped	3
4	10159	49	100.00	14	5205.27	10/10/2003 0:00	Shipped	4

5 rows × 25 columns

```
In [5]: df.shape
```

```
Out[5]: (2823, 25)
```

```
In [6]: df.describe()
```

```
Out[6]:
```

	ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	QTR_ID	MONTH
count	2823.000000	2823.000000	2823.000000	2823.000000	2823.000000	2823.000000	2823.000000
mean	10258.725115	35.092809	83.658544	6.466171	3553.889072	2.717676	7.0924

	ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	QTR_ID	MONTH
std	92.085478	9.741443	20.174277	4.225841	1841.865106	1.203878	3.6566
min	10100.000000	6.000000	26.880000	1.000000	482.130000	1.000000	1.0000
25%	10180.000000	27.000000	68.860000	3.000000	2203.430000	2.000000	4.0000
50%	10262.000000	35.000000	95.700000	6.000000	3184.800000	3.000000	8.0000
75%	10333.500000	43.000000	100.000000	9.000000	4508.000000	4.000000	11.0000
max	10425.000000	97.000000	100.000000	18.000000	14082.800000	4.000000	12.0000

In [7]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2823 entries, 0 to 2822
Data columns (total 25 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ORDERNUMBER           2823 non-null   int64
1   QUANTITYORDERED       2823 non-null   int64
2   PRICEEACH             2823 non-null   float64
3   ORDERLINENUMBER       2823 non-null   int64
4   SALES                 2823 non-null   float64
5   ORDERDATE             2823 non-null   object
6   STATUS                2823 non-null   object
7   QTR_ID                2823 non-null   int64
8   MONTH_ID              2823 non-null   int64
9   YEAR_ID               2823 non-null   int64
10  PRODUCTLINE           2823 non-null   object
11  MSRP                  2823 non-null   int64
12  PRODUCTCODE           2823 non-null   object
13  CUSTOMERNAME          2823 non-null   object
14  PHONE                 2823 non-null   object
15  ADDRESSLINE1           2823 non-null   object
16  ADDRESSLINE2           302 non-null    object
17  CITY                  2823 non-null   object
18  STATE                 1337 non-null   object
19  POSTALCODE            2747 non-null   object
20  COUNTRY               2823 non-null   object
21  TERRITORY             1749 non-null   object
22  CONTACTLASTNAME        2823 non-null   object
23  CONTACTFIRSTNAME       2823 non-null   object
24  DEALSIZE              2823 non-null   object
dtypes: float64(2), int64(7), object(16)
memory usage: 551.5+ KB
```

In [8]: `df.isnull().sum()`

```
Out[8]: ORDERNUMBER           0
QUANTITYORDERED         0
PRICEEACH                0
ORDERLINENUMBER          0
SALES                    0
ORDERDATE                0
STATUS                   0
QTR_ID                   0
MONTH_ID                 0
YEAR_ID                  0
PRODUCTLINE              0
MSRP                     0
PRODUCTCODE              0
```

```
CUSTOMERNAME      0
PHONE              0
ADDRESSLINE1       0
ADDRESSLINE2       2521
CITY               0
STATE             1486
POSTALCODE         76
COUNTRY            0
TERRITORY          1074
CONTACTLASTNAME    0
CONTACTFIRSTNAME   0
DEALSIZE           0
dtype: int64
```

```
In [9]: df.dtypes
```

```
Out[9]: ORDERNUMBER      int64
QUANTITYORDERED    int64
PRICEEACH          float64
ORDERLINENUMBER    int64
SALES              float64
ORDERDATE          object
STATUS             object
QTR_ID            int64
MONTH_ID          int64
YEAR_ID           int64
PRODUCTLINE        object
MSRP              int64
PRODUCTCODE        object
CUSTOMERNAME       object
PHONE              object
ADDRESSLINE1       object
ADDRESSLINE2       object
CITY               object
STATE             object
POSTALCODE         object
COUNTRY            object
TERRITORY          object
CONTACTLASTNAME    object
CONTACTFIRSTNAME   object
DEALSIZE           object
dtype: object
```

```
In [10]: df_drop = ['ADDRESSLINE1', 'ADDRESSLINE2', 'STATUS', 'POSTALCODE', 'CITY', 'TERRITORY', 'P
df = df.drop(df_drop, axis=1) #Dropping the categorical unnecessary columns along with co.
```

```
In [11]: df.isnull().sum()
```

```
Out[11]: QUANTITYORDERED    0
PRICEEACH                  0
ORDERLINENUMBER           0
SALES                     0
ORDERDATE                 0
QTR_ID                   0
MONTH_ID                 0
YEAR_ID                  0
PRODUCTLINE              0
MSRP                    0
PRODUCTCODE              0
COUNTRY                  0
DEALSIZE                 0
dtype: int64
```

```
In [12]: df.dtypes
```

```
Out[12]: QUANTITYORDERED      int64
PRICEEACH                    float64
ORDERLINENUMBER              int64
SALES                        float64
ORDERDATE                    object
QTR_ID                       int64
MONTH_ID                     int64
YEAR_ID                      int64
PRODUCTLINE                  object
MSRP                         int64
PRODUCTCODE                  object
COUNTRY                      object
DEALSIZE                     object
dtype: object
```

```
In [13]: # Checking the categorical columns.
```

```
In [14]: df['COUNTRY'].unique()
```

```
Out[14]: array(['USA', 'France', 'Norway', 'Australia', 'Finland', 'Austria', 'UK',
        'Spain', 'Sweden', 'Singapore', 'Canada', 'Japan', 'Italy',
        'Denmark', 'Belgium', 'Philippines', 'Germany', 'Switzerland',
        'Ireland'], dtype=object)
```

```
In [15]: df['PRODUCTLINE'].unique()
```

```
Out[15]: array(['Motorcycles', 'Classic Cars', 'Trucks and Buses', 'Vintage Cars',
        'Planes', 'Ships', 'Trains'], dtype=object)
```

```
In [16]: df['DEALSIZE'].unique()
```

```
Out[16]: array(['Small', 'Medium', 'Large'], dtype=object)
```

```
In [17]: productline = pd.get_dummies(df['PRODUCTLINE']) #Converting the categorical columns.
Dealsize = pd.get_dummies(df['DEALSIZE'])
```

```
In [18]: df = pd.concat([df,productline,Dealsize], axis = 1)
```

```
In [19]: df_drop = ['COUNTRY','PRODUCTLINE','DEALSIZE'] #Dropping Country too as there are alot of
df = df.drop(df_drop, axis=1)
```

```
In [20]: df['PRODUCTCODE'] = pd.Categorical(df['PRODUCTCODE']).codes #Converting the datatype.
```

```
In [21]: df.drop('ORDERDATE', axis=1, inplace=True) #Dropping the Orderdate as Month is already included
```

```
In [22]: df.dtypes #All the datatypes are converted into numeric
```

```
Out[22]: QUANTITYORDERED      int64
PRICEEACH                    float64
ORDERLINENUMBER              int64
SALES                        float64
QTR_ID                       int64
```

```

MONTH_ID      int64
YEAR_ID       int64
MSRP           int64
PRODUCTCODE   int8
Classic Cars  uint8
Motorcycles   uint8
Planes        uint8
Ships         uint8
Trains        uint8
Trucks and Buses uint8
Vintage Cars  uint8
Large         uint8
Medium        uint8
Small         uint8
dtype: object

```

Plotting the Elbow Plot to determine the number of clusters.

```


In [23]: distortions = [] # Within Cluster Sum of Squares from the centroid
K = range(1,10)
for k in K:
    kmeanModel = KMeans(n_clusters=k)
    kmeanModel.fit(df)
    distortions.append(kmeanModel.inertia_) #Appending the inertia to the Distortions

```

```

In [24]: plt.figure(figsize=(16,8))
plt.plot(K, distortions, 'bx-')
plt.xlabel('k')
plt.ylabel('Distortion')
plt.title('The Elbow Method showing the optimal k')
plt.show()

```



As the number of k increases Inertia decreases.

Observations: A Elbow can be observed at 3 and after that the curve decreases gradually.

```
In [25]: X_train = df.values #Returns a numpy array.
```

```
In [26]: X_train.shape
```

```
Out[26]: (2823, 19)
```

```
In [27]: model = KMeans(n_clusters=3,random_state=2) #Number of cluster = 3  
model = model.fit(X_train) #Fitting the values to create a model.  
predictions = model.predict(X_train) #Predicting the cluster values (0,1,or 2)
```

```
In [28]: unique,counts = np.unique(predictions,return_counts=True)
```

```
In [29]: counts = counts.reshape(1,3)
```

```
In [30]: counts_df = pd.DataFrame(counts,columns=['Cluster1','Cluster2','Cluster3'])
```

```
In [31]: counts_df.head()
```

```
Out[31]:
```

	Cluster1	Cluster2	Cluster3
0	1083	1367	373

Visualization

```
In [32]: pca = PCA(n_components=2) #Converting all the features into 2 columns to make it easy to v
```

```
In [33]: reduced_X = pd.DataFrame(pca.fit_transform(X_train),columns=['PCA1','PCA2']) #Creating a l
```


```
In [34]: reduced_X.head()
```

```
Out[34]:
```

	PCA1	PCA2
0	-682.488323	-42.819535
1	-787.665502	-41.694991
2	330.732170	-26.481208
3	193.040232	-26.285766
4	1651.532874	-6.891196

```
In [35]: #Plotting the normal Scatter Plot  
plt.figure(figsize=(14,10))  
plt.scatter(reduced_X['PCA1'],reduced_X['PCA2'])
```

Out[35]: <matplotlib.collections.PathCollection at 0x23b761b17f0>



In [36]: `model.cluster_centers_` *#Finding the centriods. (3 Centriods in total. Each Array contains*

Out[36]: `array([[3.72031394e+01, 9.52120960e+01, 6.44967682e+00,`

In [37]: `reduced_centers = pca.transform(model.cluster_centers_)` *#Transforming the centroids into*


In [38]: `reduced_centers`

Out[38]: `array([[5.84994044e+02, -4.36786931e+00],`

```
[ -1.43005891e+03,  2.60041009e+00],  
[ 3.54247180e+03,  3.15185487e+00]])
```

```
In [39]: plt.figure(figsize=(14,10))  
plt.scatter(reduced_X['PCA1'],reduced_X['PCA2'])  
plt.scatter(reduced_centers[:,0],reduced_centers[:,1],color='black',marker='x',s=300) #Plot
```

```
Out[39]: <matplotlib.collections.PathCollection at 0x23b75cedca0>
```



```
In [40]: reduced_X['Clusters'] = predictions #Adding the Clusters to the reduced dataframe.
```

```
In [41]: reduced_X.head()
```

```
Out[41]:
```

	PCA1	PCA2	Clusters
0	-682.488323	-42.819535	1
1	-787.665502	-41.694991	1
2	330.732170	-26.481208	0
3	193.040232	-26.285766	0
4	1651.532874	-6.891196	0


```
In [42]: #Plotting the clusters  
plt.figure(figsize=(14,10))  
# taking the cluster number and first column taking the same  
plt.scatter(reduced_X[reduced_X['Clusters'] == 0].loc[:, 'PCA1'],reduced_X[reduced_X['Clusters'] == 0].loc[:, 'PCA2'])  
plt.scatter(reduced_X[reduced_X['Clusters'] == 1].loc[:, 'PCA1'],reduced_X[reduced_X['Clusters'] == 1].loc[:, 'PCA2'])
```



```
plt.scatter(reduced_X[reduced_X['Clusters'] == 2].loc[:, 'PCA1'], reduced_X[reduced_X['Clusters'] == 2].loc[:, 'PCA2'], color='red', marker='x', s=300)
```

```
plt.scatter(reduced_centers[:,0], reduced_centers[:,1], color='black', marker='x', s=300)
```

Out[42]: <matplotlib.collections.PathCollection at 0x23b75d35a00>



In []: